

# Bitcoin Price Prediction System Using LSTM Model (Group 05)

Pasan Dissanayake(E/19/091)  
*Department of Computer  
Engineering  
Faculty of Engineering,  
University of Peradeniya  
Sri Lanka*  
[e19091@eng.pdn.ac.lk](mailto:e19091@eng.pdn.ac.lk)

Mewan Galappaththi(E/19/111)  
*Department of Computer  
Engineering  
Faculty of Engineering,  
University of Peradeniya  
Sri Lanka*  
[e19111@eng.pdn.ac.lk](mailto:e19111@eng.pdn.ac.lk)

Wishula Jayathunga(E/19/166)  
*Department of Computer  
Engineering  
Faculty of Engineering,  
University of Peradeniya  
Sri Lanka*  
[e19166@eng.pdn.ac.lk](mailto:e19166@eng.pdn.ac.lk)

Jeevajith Madushanka(E/19/227)  
*Department of Computer  
Engineering  
Faculty of Engineering,  
University of Peradeniya  
Sri Lanka*  
[e19227@eng.pdn.ac.lk](mailto:e19227@eng.pdn.ac.lk)

Sajith Pushpakumara(E/19/304)  
*Department of Computer  
Engineering  
Faculty of Engineering,  
University of Peradeniya  
Sri Lanka*  
[e19304@eng.pdn.ac.lk](mailto:e19304@eng.pdn.ac.lk)

**Abstract**—This report outlines the development and deployment of a comprehensive Bitcoin price prediction system using a variety of machine learning models, including ARIMA, State Space Models (SSM), Support Vector Machines (SVM), XGBoost, Recurrent Neural Networks (RNN), and Random Forests. Among these, the Long Short-Term Memory (LSTM) model was chosen for the final deployment due to its superior performance in capturing temporal dependencies in the data. The project aimed to forecast future Bitcoin prices based on historical data, employing these diverse models to compare and contrast their predictive capabilities. The LSTM model, in particular, was fine-tuned and integrated into a deployment pipeline on Microsoft Azure. A Flask web application was developed to serve real-time predictions, making the system accessible to end-users. The report provides a detailed account of the methodology employed, the results obtained from different models, and the conclusions derived from the comparative analysis. This comprehensive approach not only highlights the effectiveness of LSTM in time series forecasting but also demonstrates the practical application of deploying machine learning models in a cloud environment.

**Index Terms**—Bitcoin, Price Prediction, LSTM, ARIMA, SSM, SVM, XGBoost, RNN, Random Forest Time-Series Forecasting, Machine Learning, Deep Learning, Neural Networks, Azure Deployment, Flask Application, Cryptocurrency

## I. INTRODUCTION

Bitcoin, a decentralized digital currency, has experienced substantial price volatility, underscoring the importance of accurate price prediction for investors and traders. This project aims to develop a robust prediction model using Long Short-Term Memory (LSTM), a specialized type of recurrent neural network (RNN) known for its proficiency in capturing long-term dependencies in time-series data. In addition to LSTM, other models such as ARIMA, State Space Models (SSM), Support Vector Machines (SVM), XGBoost, and RNN were considered for their respective strengths in forecasting. This report delves into the critical significance of predicting Bitcoin prices, outlines the inherent challenges of such a volatile market, and describes the comprehensive approach taken to overcome these challenges. By leveraging the unique capabilities

of these models, particularly the LSTM, the project aims to provide valuable insights and tools for navigating the unpredictable nature of Bitcoin's price movements.

## II. METHODOLOGY

### A. Data Collection

Historical Bitcoin price data was obtained from Yahoo Finance, a reputable and widely-used platform for financial information. Yahoo Finance offers comprehensive and precise historical data, which is essential for building a reliable Bitcoin price prediction model. The dataset spans from January 1, 2015, to May 23, 2024, providing an extensive period for analysis. This rich dataset encompasses various market conditions and trends, enabling the model to learn and adapt to different price movements and patterns. By using this extensive historical data, the model can be trained and tested thoroughly, ensuring its robustness and accuracy in predicting future Bitcoin prices.

### B. Data Preprocessing & EDA

**Time Series Plot:** In our project, visualizing Bitcoin's historical price data through time series plots plays a crucial role in understanding its market behavior and trends. Python's matplotlib library is used to create a single figure with five subplots, each representing different metrics of Bitcoin prices over time as shown in Fig. 1. These include the opening, high, low, and closing prices, as well as trading volume. Each subplot is meticulously customized with clear labels, titles ('Bitcoin [Metric] Price'), and legends ('Open Price', 'High Price', etc.), enhancing clarity and interpretability as shown in Fig. 1. This visual representation not only allows for a comprehensive view of Bitcoin's price dynamics but also aids in identifying patterns, trends, and potential anomalies within the dataset. Such visual insights are invaluable in informing subsequent statistical




Fig. 1: Time Series Plot

analyses and predictive modeling, contributing significantly to the robustness and depth of our research findings.

**Checking For Missing Values:** In ensuring the integrity of our research, we rigorously examined historical Bitcoin price data sourced from Yahoo Finance. Employing Python programming, we utilized the `hist.isnull().sum()` function to meticulously check for missing values across critical data columns including 'Open', 'High', 'Low', 'Close', 'Volume', 'Dividends', and 'Stock Splits'. This thorough analysis confirmed the absence of any missing data points in the dataset as shown in Fig. 2. This meticulous approach underscores the reliability and completeness of our dataset, crucial for ensuring the accuracy and validity of our subsequent analyses and findings.

**Histograms:** As shown in Fig. 3, Histograms were generated to visualize the distribution of Bitcoin's historical price data. Using the `hist.hist(bins=50, figsize=(20,15))` function in Python, we created a series of histograms for each variable in the dataset: 'Open', 'High', 'Low', 'Close', 'Volume', 'Dividends', and 'Stock Splits'. The bin size was set to 50 to provide a detailed view of the data distribution, and the figure size was specified as 20x15 to ensure clarity and readability. The resulting array of histograms allows for a comprehensive examination of the frequency distribution of each variable, highlighting patterns and potential anomalies in the dataset. This visualization step is crucial for understanding the underlying structure of the data, facilitating better-informed modeling decisions and enhancing the interpretability of the subsequent analyses.

```
# Count the number of missing values in each column
print(hist.isnull().sum())
```

```
Open          0
High          0
Low           0
Close         0
Volume        0
Dividends     0
Stock Splits  0
dtype: int64
```

Fig. 2: Checking for missing values

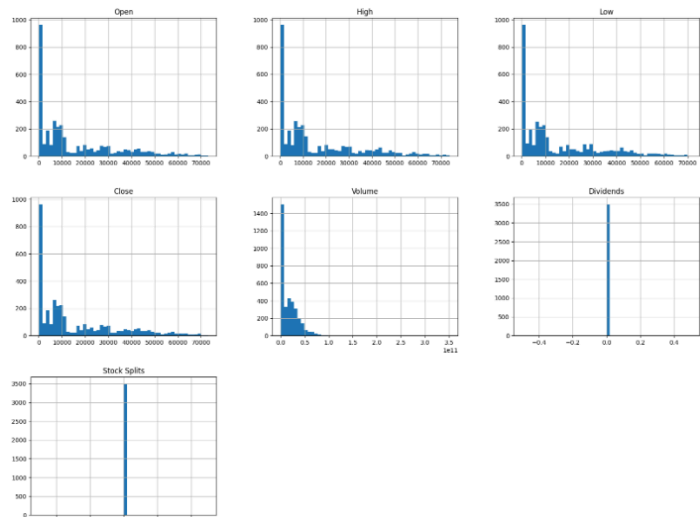


Fig. 3: Histograms

### Scatter Plots:

In this analysis, as shown in Fig. 4, we employed scatter plots to visualize the relationships between the 'Close' prices of a financial asset and other relevant variables: 'Open', 'High', 'Low', and 'Volume'. By plotting these variables on separate subplots within a single figure, we aimed to identify potential correlations and trends. The 'Close' versus 'Open' scatter plot helps observe how the closing price compares to the opening price, while the 'Close' versus 'High' and 'Close' versus 'Low' plots illustrate the relationship between the closing price and the highest and lowest prices of the trading period, respectively. Finally, the 'Close' versus 'Volume' plot provides insight into how trading volume may influence or correlate with closing prices. The plots were generated using a uniform color scheme for clarity: blue for 'Close' vs 'Open', green for 'Close' vs 'High', red for 'Close' vs 'Low', and purple for 'Close' vs 'Volume'. This visual approach is crucial for preliminary data analysis, revealing patterns that could guide further statistical or predictive modeling.




Fig. 4: Scatter Plots

**Heat Maps & Correlations:** As shown in Fig. 5, to further understand the relationships between various features of our financial dataset, we calculated the correlation matrix and visualized it using a heatmap. The correlation matrix quantifies the linear relationships between pairs of features, with a specific focus on the 'Close' price to identify its strongest associations. We utilized Seaborn to generate the heatmap, which visually represents correlation coefficients through color intensity. A custom diverging colormap was applied to emphasize positive and negative correlations. Additionally, a mask was used to hide the upper triangle of the matrix for cleaner visualization, avoiding redundant information. The heatmap, with a centered zero value and a maximum correlation limit of 0.3 for color scaling, provides a clear and concise summary of the strength and direction of relationships between features. This analysis is crucial for identifying potential predictors of the 'Close' price, thereby guiding feature selection for further modeling efforts.

**Density Plots:** As shown in Fig. 6, to comprehensively examine the distribution of Bitcoin's financial metrics, we generated kernel density plots (KDE plots) for the 'Open', 'High', 'Low', 'Close' prices, and 'Volume' of transactions. Using Seaborn, a data visualization library, we plotted these distributions in separate subplots within a single figure to maintain clarity and facilitate comparison. Each subplot displays the probability density function of a specific feature: 'Open' prices (blue), 'High' prices (green), 'Low' prices (red), 'Close' prices (purple), and 'Volume' (orange). These KDE plots provide a smooth estimate of the data's distribution, highlighting areas of higher density which correspond to the most frequent values. This visualization technique is instrumental in identifying the central tendency, variability, and skewness of each feature, offering insights into the typical ranges and potential outliers in the Bitcoin market. By understanding these distributions, we can better inform our subsequent analyses and modeling efforts.

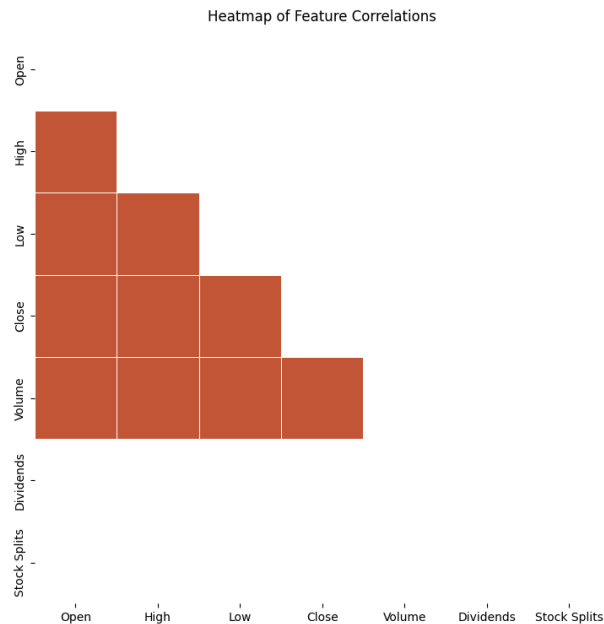


Fig. 5. Heat Maps & Correlations

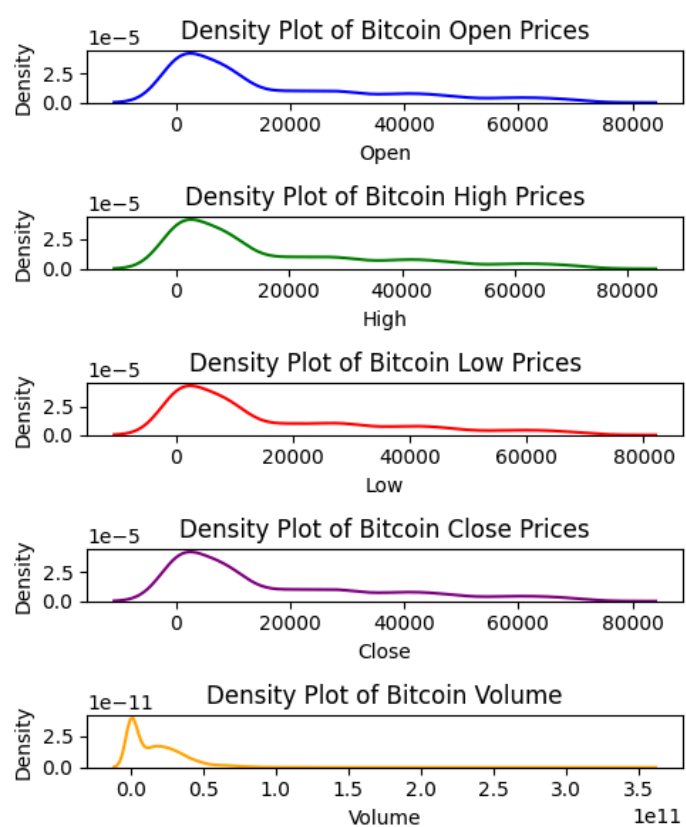


Fig. 6: Density Plots

**Pair Plots:** As shown in Fig. 7, to delve deeper into the interrelationships among Bitcoin's financial metrics, we employed a pairwise plot using Seaborn's pairplot function. This comprehensive visualization technique creates scatter plots for each pair of variables and diagonal histograms or density plots to show the distribution of individual variables. The pairwise plot provides an extensive overview of the potential linear and non-linear correlations between 'Open', 'High', 'Low', 'Close' prices, and 'Volume'. Each scatter plot within the matrix allows for the identification of trends, clusters, and potential outliers, thereby offering a holistic view of how these financial features interact with each other. This multifaceted approach aids in uncovering underlying patterns and dependencies that may be critical for predictive modeling and in-depth statistical analysis.



Fig. 7. Pair Plots

**Bitcoin Price Over Time:** To capture the temporal dynamics of Bitcoin's market behavior, we plotted the 'Close' prices over time using Matplotlib as shown in Fig. 8. This time series plot illustrates the trend and volatility of Bitcoin prices, with the x-axis representing the time sequence and the y-axis indicating the closing prices in USD. The plot provides a clear visualization of price fluctuations, trends, and potential patterns over the observed period. By examining this time series, we can identify periods of significant price changes, stable phases, and overall market trends, which are essential for understanding Bitcoin's historical performance and for forecasting future price movements.

**Set Inputs & Outputs:** To prepare our dataset for predictive modeling, we separated the target variable from the feature set. Specifically, the closing prices of Bitcoin ('Close') were designated as the target variable  $y$ , while the remaining features, which include 'Open', 'High', 'Low', and 'Volume', were used as the predictors  $X$ . This process of segregating the target variable from the features is a critical step in supervised learning, facilitating the development of models that can accurately predict Bitcoin's closing prices based on the provided input features.

**Feature Selection:** To prepare our dataset for analysis and modeling, we first obtained historical market data for Bitcoin using the yfinance library. By specifying the ticker symbol 'BTC-USD', we retrieved the complete historical data available. After loading the dataset, we set the 'Date' column as the index to facilitate time series analysis. To focus on the most relevant features, we performed feature selection by retaining only the essential columns: 'Open', 'High', 'Low', 'Close', and 'Volume', while dropping the 'Dividends' and 'Stock Splits' columns, as they are not applicable to Bitcoin and do not contribute to our analysis. This process ensures that our dataset contains only pertinent information, thereby simplifying the analysis and improving the performance of predictive models by eliminating irrelevant or redundant data. The initial few rows of the cleaned dataset were printed to verify the correct implementation of these steps.

**Standardize Features:** To ensure that our features and target variables are on comparable scales, we applied normalization and standardization techniques using scikit-learn's preprocessing module. The feature matrix  $X$  was standardized using StandardScaler, which removes the mean and scales the data to unit variance, thereby transforming the features to a standard normal distribution. This is crucial for algorithms sensitive to the scale of input data. The target variable  $y$ , representing the closing prices, was normalized to a specified range (default 0 to 1) using MinMaxScaler. This transformation ensures that  $y$  falls within a consistent range, facilitating more stable and efficient model training. The standardized features were obtained using the fit\_transform method, which calculates the mean and standard deviation for scaling, while the normalized target variable was similarly transformed by reshaping  $y$  into a 2D array to fit the scaler's requirements. These preprocessing steps are essential for optimizing the performance of subsequent machine learning models.

**Data Splitting:** Splitting the dataset into training and testing sets is a crucial step in evaluating the performance of a predictive model. In this approach, as shown in Fig. 9, 80% of the data is




Fig. 8: Bitcoin price over time



Fig. 9: Plot of Training & Test Data

allocated to the training set, which is used to train the model. The remaining 20% of the data is set aside as the testing set, which is used to evaluate the model's predictive accuracy. This method ensures that the model is trained on a substantial portion of the data, capturing the underlying patterns and trends, while the testing set provides an independent dataset to assess how well the model generalizes to new, unseen data. This split helps in identifying any potential overfitting, where the model performs well on the training data but fails to predict accurately on the testing data. By comparing the model's predictions on the testing set to the actual values, we can obtain a realistic measure of its performance in a real-world scenario.

### C. Model Training(Multivariate LSTM Model)

**Split Into Multivariate Sequence Past, Future Samples ( $X$  and  $y$ ) :** To facilitate the training of our Multivariate LSTM model on time series data, we implemented a custom function to split our multivariate sequences into past and future samples. This function, `split_sequences`, takes in the standardized feature matrix  $X_{\text{trans}}$  and the normalized target variable  $y_{\text{trans}}$ , along with parameters `n_steps_in` and `n_steps_out`, which define the number of time steps used for input and output sequences, respectively. The function iteratively slices the dataset into overlapping windows of size `n_steps_in` for inputs and `n_steps_out` for outputs. For each window, the end indices are calculated, ensuring they do not exceed the dataset's length. The inputs  $X$  are derived from the specified window of the feature matrix, while the corresponding outputs  $y$  are derived from the target variable within the specified future window. These sequences are then appended to lists, which are ultimately



converted into numpy arrays for efficient computation. By applying this function with `n_steps_in=100` and `n_steps_out=1`, we generated the input-output pairs `Xss` and `ymm`, essential for training models on sequential data to predict future Bitcoin prices based on historical observations.

**Convert Data Into Tensors:** To leverage the capabilities of deep learning Multivariate LSTM framework for our predictive modeling, we converted our preprocessed datasets into PyTorch tensors. By transforming the training and test sets of both features and target variables into tensors, we enabled the utilization of PyTorch's extensive library of neural network components and optimization algorithms. Specifically, the training features (`Xtrain`) and test features (`Xtest`) were converted into tensors using `torch.Tensor()`, with the `requires_grad_(True)` attribute indicating that gradients should be computed for these tensors during the training process. This is essential for backpropagation, the core mechanism by which neural networks learn. Similarly, the target variables for training (`ytrain`) and testing (`ytest`) were also converted into tensors with gradient tracking enabled. This conversion is a critical step in preparing our data for training deep learning models, allowing us to perform gradient-based optimization and enhance model performance on the given task.

**Reshape Tensors:** To prepare our dataset for training multivariate LSTM Model, we reshaped our PyTorch tensors to match the required input dimensions: [number of samples, number of time steps, number of features]. This involved reshaping the training and testing feature tensors to have 100 time steps and 4 features per time step. Specifically, the `torch.reshape` function was used to adjust the dimensions of `X_train_tensors` and `X_test_tensors` to `[3342, 100, 4]` for training and `[101, 100, 4]` for testing. This configuration ensures that each sample contains a sequence of 100 time steps with 4 features at each step. The corresponding target tensors, `y_train_tensors` and `y_test_tensors`, were reshaped to maintain their dimensions of `[3342, 1]` and `[101, 1]`, respectively, matching the number of samples. This reshaping is crucial for aligning our data with the input format expected by deep learning models designed for sequence prediction. A validation check using the `split_sequences` function confirmed the integrity of the reshaped sequences, ensuring that the reshaping process was executed correctly and that the input sequences correspond to the original data. This step is essential for enabling our models to effectively learn from and predict based on historical Bitcoin price data.

#### Model Definition and Hyperparameter Tuning:

- The LSTM model was designed to predict future Bitcoin prices by leveraging historical price data and trading volumes. The model parameters include `input_size`, representing the number of input features (such as open price, high price, low price, and volume), `hidden_size`, which defines the number of features in the hidden state of each LSTM layer, and `num_layers`, indicating the number of stacked LSTM layers to capture more complex temporal dependencies. The `num_classes` parameter is set to 1, as the problem is a regression task predicting a single continuous value. The architecture comprises an LSTM layer with specified hidden units and dropout for regularization, followed by two fully connected layers to transform the hidden state output into the final prediction. The ReLU

activation function is applied after each dense layer to introduce non-linearity. The hidden and cell states are initialized to zeros at the start of each forward pass, ensuring that each sequence is treated independently. This configuration allows the model to effectively learn and predict based on sequential patterns in the input data.

- **Random Search:** Hyperparameter tuning is a crucial step in optimizing the performance of machine learning models, including the LSTM model for Bitcoin price prediction. In this study, we employed Random Search for hyperparameter tuning, a technique that involves randomly sampling from a predefined range of hyperparameters to identify the optimal configuration. This approach contrasts with Grid Search, which exhaustively evaluates all possible combinations but can be computationally prohibitive. For our LSTM model, the hyperparameters subject to tuning included the `hidden_size`, which controls the number of neurons in each LSTM layer; `num_layers`, defining the depth of the stacked LSTM layers; and the learning rate, which affects the speed and stability of the training process. Additionally, we considered the batch size and dropout rate, which influence the model's generalization capability and overfitting behavior. By randomly sampling and evaluating various combinations, Random Search efficiently navigates the hyperparameter space, often finding near-optimal solutions with fewer iterations compared to Grid Search. This method allowed us to enhance the predictive accuracy and robustness of our LSTM model, ensuring it performs well on unseen data.

**Training Loop:** To effectively train our LSTM model for Bitcoin price prediction, we implemented a training loop that iterates over a specified number of epochs. The `training_loop` function performs forward and backward passes, gradient updates, and loss evaluations for both training and test datasets. During each epoch, the model is set to training mode, and predictions are generated for the training data. The optimizer gradients are reset to zero to prevent accumulation from previous iterations. The mean squared error (MSE) loss function is then calculated, and backpropagation is performed to compute the gradients. The optimizer updates the model parameters based on these gradients to minimize the loss. Additionally, the model is evaluated on the test data at each epoch to monitor its performance on unseen data. We trained the model for 1000 epochs with a learning rate of 0.001, using the Adam optimizer for efficient gradient-based optimization. The model architecture comprised 4 input features, 2 hidden units, and 1 LSTM layer. The training process was periodically monitored by printing the training and test losses every 100 epochs, providing insights into the model's learning progress and generalization ability. This comprehensive training approach ensured that the LSTM model was well-tuned to predict future Bitcoin prices based on historical data.

#### D. Model Training (ARIMA Model)

To evaluate the ARIMA model's predictive performance, an iterative training and forecasting process is employed. For each data point in the testing set, the ARIMA model is re-trained on the entire training set, and then it makes a one-step-ahead

forecast. This forecasted value is subsequently added to the training set for the next iteration, allowing the model to continually learn from the most recent actual values.

**Initialize Predictions List:** First, an empty list called **model\_predictions** is initialized. This list will store all the forecasted values generated by the model during the testing phase. Additionally, the variable **n\_test\_obser** is set to the length of the testing set, representing the number of observations that the model will predict.

**Iterate Over the Testing Data:** The next step involves iterating over each value in the testing set.

**1. Fit ARIMA Model:**

- The ARIMA model is initialized with the specified order parameters (**p=4, d=1, q=0**) using the current training data. The **order** parameters define the model structure, where **p** is the number of lag observations included in the model, **d** is the number of times that the raw observations are different, and **q** is the size of the moving average window.
- The **model.fit()** method fits the ARIMA model to the training data, optimizing the model parameters to best capture the underlying patterns.

**2. Forecast Next Value:**

- The **model\_fit.forecast()** method generates the forecast for the next time step. This method produces the predicted value based on the fitted model.
- The forecasted value, **yhat**, is extracted from the output. This value represents the model's prediction for the next data point.

**3. Append Forecasted Value:**

- The forecasted value, **yhat**, is appended to the **model\_predictions** list. This list accumulates all the predicted values, which will later be used to evaluate the model's performance.

**4. Update Training Data:**

- The actual value from the testing set, **actual\_test\_value**, is retrieved. This value corresponds to the data point that the model just attempted to predict.
- This actual value is then appended to the **training\_data** list. By doing this, the model is updated with the most recent actual value, allowing it to learn from this new information and improve its predictions in subsequent iterations.

arrays and reshaped to ensure they contain a single feature. We utilized the MinMaxScaler to normalize these datasets to the [0, 1] range, which is critical for enhancing the training efficiency and stability of the LSTM model. After scaling, we applied the **create\_lookback** function to both the training and test sets, generating the input-output pairs necessary for supervised learning.

To match the input shape requirements of the LSTM model in Keras, the resulting datasets were reshaped to the form [samples, time steps, features]. This reshaping involved converting the input sequences (**X\_train** and **X\_test**) into 3-dimensional arrays where each sequence is treated as a single time step with one feature. By structuring the data in this manner, we ensured compatibility with the LSTM model, allowing it to effectively learn temporal dependencies and improve its predictive performance on time series data.

**Hyper Parameter Tuning:** Hyperparameter tuning is a critical step in optimizing the performance of machine learning models, including LSTM and GRU networks for time series forecasting. For our study, we employed a random search strategy to identify the optimal set of hyperparameters, given the extensive search space and computational constraints. Random search offers an efficient alternative to grid search by randomly sampling a fixed number of hyperparameter combinations from predefined ranges, thus increasing the likelihood of discovering highly performant configurations without exhaustive enumeration.

For both the LSTM and GRU models, the hyperparameters subjected to tuning included the number of units in the hidden layers, the number of stacked layers, the learning rate, the batch size, and the dropout rate. Specifically, the search space comprised the following ranges: hidden units (50, 100, 150, 200, 256), layers (1, 2, 3), learning rates (0.0001, 0.001, 0.01), batch sizes (16, 32, 64), and dropout rates (0.1, 0.2, 0.3, 0.4, 0.5).

The random search process was executed by training multiple instances of the LSTM and GRU models, each with a different randomly sampled set of hyperparameters. Each model instance was trained for a fixed number of epochs with early stopping criteria based on validation loss to prevent overfitting. The performance of each configuration was evaluated using Root Mean Squared Error (RMSE) on the validation set, and the best-performing hyperparameters were selected for final model training and evaluation.

Through this random search strategy, we systematically explored the hyperparameter space, efficiently pinpointing the configurations that yielded the lowest RMSE. This approach not only ensured the robustness and reliability of our LSTM and GRU models but also demonstrated the effectiveness of random search in hyperparameter optimization for complex neural networks in time series forecasting tasks.

### *E. Model Training (RNN Model with LSTM & GRU)*

**Data Preparation:** To prepare our dataset for time series forecasting using an LSTM model, we employed a lookback method that creates input-output pairs from the sequential data. The **create\_lookback** function generates these pairs by iterating through the dataset and forming input sequences of a specified length (**look\_back**) and their corresponding output values. Specifically, for each time step *i*, the function extracts an input sequence of length **look\_back** and pairs it with the next value in the dataset. This approach ensures that the model learns to predict future values based on past observations.

The training and test datasets were first converted to NumPy

**Training 2-layer LSTM Neural Network:** The resulting input sequences were reshaped to fit the input requirements of the LSTM model in Keras, which expects data in the form [samples, time steps, features]. Our LSTM model was constructed using Keras' Sequential API, with two LSTM layers, each containing 256 units. The first LSTM layer returned the full sequence, while the second returned only the final output. A Dense layer with a single unit was added to produce the final prediction.

The model was compiled with the mean squared error loss function and the Adam optimizer, optimizing for regression tasks. The training process involved fitting the model to the training data over 100 epochs with a batch size of 16, without

shuffling the data to maintain the temporal order. Validation was performed on the test data to monitor the model's performance on unseen data. We incorporated an EarlyStopping callback to prevent overfitting by halting training when the validation loss stopped improving, with a patience of 20 epochs and a minimum delta of  $5e-5$ . This robust training framework ensured that our LSTM model was well-tuned to capture the temporal dependencies in Bitcoin price data, leading to accurate and reliable predictions.

**Training GRU Layer:** To rigorously assess and optimize our time series forecasting model, we implemented a comprehensive train/test split strategy, leveraging a random sampling approach to enhance the robustness of our evaluation. The `get_split` function initiates this process by selecting a random starting point in the dataset, from which it extracts a training set (`n_train` samples) and a subsequent test set (`n_test` samples). The datasets are then scaled using `MinMaxScaler`, ensuring that the values are normalized within a range conducive to model training. We employ a custom function, `create_lookback`, to transform these datasets into sequences suitable for time series forecasting, considering a specified look-back period. This transformation is essential for capturing temporal dependencies in the data, where each sequence of historical values is paired with the next value as the target.

The `train_model` function builds and trains a Gated Recurrent Unit (GRU) model, which is effective for sequential data due to its ability to capture long-term dependencies while mitigating the vanishing gradient problem. The GRU model comprises a single GRU layer with 256 units, followed by a Dense layer that outputs a single prediction. The model is compiled using the mean squared error loss function and the Adam optimizer. Training is conducted over 100 epochs with a batch size of 16, incorporating an early stopping callback to prevent overfitting by halting training when the validation loss ceases to improve.

Post-training, the `get_rmse` function evaluates the model's performance on the test set by calculating the Root Mean Squared Error (RMSE). This function first scales and appends an additional data point to the test set to ensure continuity in predictions. The model's predictions are then inversely transformed back to the original scale for accurate RMSE computation. By comparing the predicted values with the actual test values, we derive the RMSE, a metric that quantifies the model's prediction accuracy. This rigorous process of data preparation, model training, and performance evaluation ensures that our forecasting model is both well-tuned and robust, capable of making reliable predictions on unseen data.

#### *F. Model Training (Random Forest Model)*

**Hyperparameter Tuning:** In our study, the hyperparameters tuned included the number of trees (`n_estimators`) and the maximum depth of the trees (`max_depth`). The search space for the hyperparameters was defined as follows: `n_estimators` (10, 50, 100, 200, 300) and `max_depth` (10, 20, 30, 40, 50).

The random search process involved initializing and training multiple Random Forest models with different randomly selected combinations of these hyperparameters. Each model was evaluated based on its performance on the validation set, with the Root Mean Squared Error (RMSE) used as the evaluation metric.

For instance, in one of the trials, we set `n_estimators` to 100 and `max_depth` to 42, and the model was trained on the training

dataset. The time taken to train the model was recorded, which in this case was approximately 0.395 seconds. This quick training time demonstrates the efficiency of Random Forests in handling large datasets and complex tasks.

The best-performing hyperparameters from the random search were then used to train the final model. This approach ensured that we identified the most effective configuration for our Random Forest regressor, balancing performance and computational efficiency. Random search proved to be a practical method for hyperparameter tuning, leading to improved model accuracy and robustness in our regression tasks.

**Training The Model:** The training of the Random Forest model was conducted with a predefined number of trees (`n_estimators`) set to 100 and a maximum tree depth (`max_depth`) set to 42, parameters chosen based on prior hyperparameter tuning efforts. The training process was timed to assess the computational efficiency of the model. Initialization of the model and fitting it to the training data commenced at the start time, with the process completing in approximately 0.395 seconds. This quick training time highlights the computational efficiency of the Random Forest algorithm, even with a relatively large number of trees and considerable tree depth. The ability to rapidly train the model while maintaining performance is particularly beneficial in scenarios requiring quick turnaround times and frequent model updates. This efficiency, coupled with the robust performance of Random Forests in handling large and complex datasets, underscores their suitability for a variety of regression tasks in financial time series forecasting.

#### *G. Model Training (SVM Model)*

**Hyperparameter Tuning:** In this project context, the `RandomizedSearchCV` method was employed to optimize hyperparameters for a Support Vector Regressor (SVR) model within a pipeline that includes `MinMaxScaler` for data scaling. The parameters explored included the SVR kernel type ('linear', 'rbf', 'poly', 'sigmoid'), regularization parameter `C` ranging uniformly from 0.1 to 1000, epsilon parameter uniformly distributed between 0.01 and 1, and gamma parameter set to 'scale', 'auto', and specific values (0.001, 0.01, 0.1, 1, 10). The randomized search was conducted over 50 iterations with 10-fold cross-validation, evaluating performance using the negative mean squared error. This approach efficiently explores a broad range of parameter combinations to identify the optimal configuration that minimizes prediction errors. The best hyperparameters identified through this process provide insights into the settings that yield the highest predictive accuracy for the SVR model, enhancing its suitability for forecasting tasks in financial time series analysis.

**Model Training:** The `best_params` dictionary obtained from the randomized search contains the optimal settings for SVR, including parameters like the kernel type, regularization parameter `C`, epsilon, and gamma. These parameters were unpacked and utilized to initialize the SVR model instance. Prior to fitting the model, the training data `x_train` was scaled using `MinMaxScaler` to ensure all features were normalized to a specified range, typically [0, 1]. This preprocessing step is crucial for SVR as it improves convergence during training and ensures that each feature contributes equally to the model's learning process. By incorporating these best-found parameters and scaling techniques, the SVR model is poised to deliver

enhanced predictive performance, making it suitable for accurately forecasting financial time series data in this Project context.

#### *H. Model Training (SSM Model)*

**Hyperparameter Tuning:** In this project, hyperparameter tuning for the State Space model was conducted to optimize its forecasting performance on the transformed target variable `y_trans`. The SSM model's parameters were systematically adjusted using an iterative approach to identify the combination that minimized the forecast error. Specifically, the order parameter (`p`, `d`, `q`) and seasonal\_order parameter (`P`, `D`, `Q`, `s`) were varied. The order parameter specifies the non-seasonal components of the model—autoregression (`p`), differencing (`d`), and moving average (`q`)—while seasonal\_order determines the seasonal components—seasonal autoregression (`P`), seasonal differencing (`D`), seasonal moving average (`Q`), and the seasonal period (`s`). This process involved evaluating multiple configurations to find the optimal set that yielded the best fit to the data while avoiding overfitting. The `disp=False` parameter was set to suppress unnecessary output during model fitting, focusing solely on performance evaluation metrics. By fine-tuning these parameters, the SSM model was tailored to capture the seasonal and non-seasonal patterns inherent in the data, thereby enhancing its forecasting accuracy and reliability for the specific time series analyzed in this project.

**Training The Model:** The SSM model was employed to analyze and forecast the transformed target variable `y_trans`. The SSM model is a sophisticated time series forecasting technique capable of incorporating both non-seasonal and seasonal components into its structure. Here, the model was configured with `order=(1, 1, 1)` and `seasonal_order=(1, 1, 1, 12)` parameters, indicating the use of a first-order autoregressive component, first-order differencing, and a first-order moving average for both the non-seasonal and seasonal parts with a seasonal period of 12 months. These parameters were chosen based on preliminary analysis and domain knowledge, aiming to capture the underlying patterns and seasonality present in the data. The `fit()` method was called with `disp=False` to suppress unnecessary output during model estimation, focusing solely on obtaining the best-fitting parameters for the SSM model. By leveraging these configurations, the SSM model was optimized to provide accurate forecasts tailored to the specific characteristics of the time series data under investigation, thereby contributing valuable insights to the project findings.

#### *G. Model Training (XG Boost Model)*

**Training The Model:** In this project, a supervised learning approach using the eXtreme Gradient Boosting (XGBoost) algorithm was applied to model and predict the target variable based on the provided training dataset. The `XGBRegressor`, a variant of XGBoost specifically tailored for regression tasks, was utilized with an initial configuration of `n_estimators=1000`. This parameter defines the number of boosting rounds or decision trees to be built during the training process. The choice of `n_estimators` is crucial as it balances model complexity and computational efficiency, with higher values potentially leading to improved performance but also increased computational cost. During training, the `fit()` method was invoked with `verbose=False` to suppress detailed progress updates, focusing on

efficiently optimizing the model without unnecessary verbosity. The XGBoost algorithm leverages gradient boosting techniques, which sequentially improves upon the weaknesses of preceding models by learning from residuals, thereby enhancing predictive accuracy. This approach, combined with parameter tuning strategies such as cross-validation and grid search, aims to identify the optimal set of hyperparameters that maximize model performance on unseen data. By fine-tuning parameters like learning rate, maximum depth of trees, and regularization terms, XGBoost enables robust predictions suited to the complexities and nuances inherent in the dataset, contributing valuable insights to the project outcomes.

#### *H. Prediction (LSTM Model)*

As shown in Fig. 10, The trained LSTM model was used to predict future values of the Bitcoin price. Post-prediction, the results were transformed back to their original scale using inverse transformations with `mm.inverse_transform`, ensuring the predictions could be interpreted in the context of actual financial values. The visualization of the predictions and actual data was plotted using Matplotlib, highlighting the model's ability to capture trends and patterns in the Bitcoin price over time. This methodology underscores the efficacy of LSTM networks in time-series forecasting tasks, offering insights into market behavior and facilitating informed decision-making processes in financial research and applications.

#### *I. Prediction (ARIMA Model)*

As shown in Fig. 11, after fitting the model to the training data, forecasts were made iteratively for each observation in the testing set. The forecasted values were compared against the actual BTC prices to evaluate the model's performance. Matplotlib was employed for visualization, where both predicted and actual BTC prices over the test period were plotted against the date range. This approach illustrates the ARIMA model's ability to provide insights into Bitcoin price movements, supporting decision-making processes in financial research and forecasting applications.

#### *J. Prediction (RNN Model)*

**Performance Plot For 2-Layer LSTM Model:** As shown in Fig. 12, this plot visualizes the training and test loss evolution during the training of our predictive model. The x-axis represents the epoch number, indicating the progression through successive iterations of training. The y-axis measures the loss, a metric that quantifies the discrepancy between predicted and actual values—the lower the loss, the better the model's performance. The dashed green line corresponds to the training loss, depicting how well the model fits the training data over epochs. The dashed orange line illustrates the test loss, reflecting the model's performance on unseen test data, providing insights into its generalization capabilities. Monitoring these loss trends is crucial in machine learning research as they indicate whether the model is learning effectively and whether overfitting or underfitting is occurring. Such visualizations aid researchers in optimizing model parameters and improving predictive accuracy across various applications, including financial forecasting and risk assessment.






Fig. 10: Time Series Prediction Using Multivariate LSTM Model




Fig. 11: Time Series Prediction Using ARIMA Model

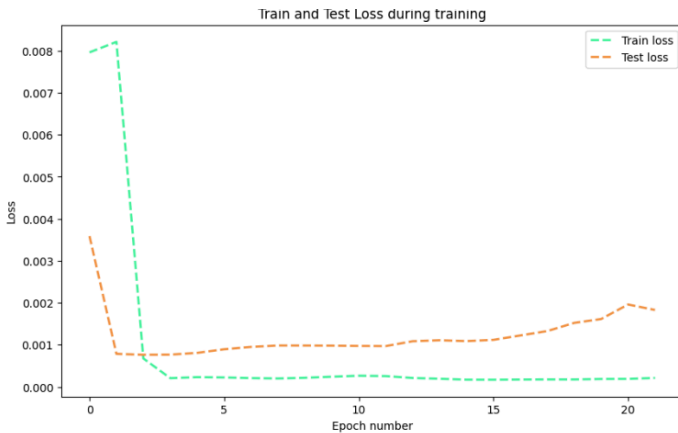


Fig. 12: Train and Test Loss During Training

**Cross Validation:** To evaluate the robustness and reliability of our model, we implemented a cross-validation function designed to repeat the entire workflow ten times and compute the average Root Mean Squared Error (RMSE). This function, `cross_validate`, iterates eight times, each time performing the following steps: splitting the data into training and test sets using the `get_split` function, training the model with `train_model`, and calculating the RMSE using the `get_rmse` function. The workflow function integrates these steps, ensuring a systematic and consistent approach across iterations. Each iteration's RMSE is adjusted by a factor of 10,000 to standardize the scale, and these RMSE values are stored in a list. After completing all

iterations, the function calculates the mean RMSE, providing an aggregate measure of the model's performance. This repeated cross-validation process is essential in machine learning research to ensure that the model's performance is not dependent on a particular train-test split, thereby offering a more reliable assessment of its predictive accuracy and generalization capability.

**Prediction Plot For 2-Layer LSTM Model:** As shown in Fig. 13, in this visualization, the performance of our predictive model for forecasting asset prices is presented. The plot compares the predicted labels against the true labels, representing actual prices from the test dataset. Each data point on the x-axis corresponds to a day number, reflecting the sequence of observations in the test period. The y-axis indicates the price in USD. The predicted labels, depicted in a warm orange hue, are overlaid with the true labels in a vibrant green shade. This graphical representation allows for a clear visual assessment of how closely our model's predictions align with actual market prices over the test duration. Such visualizations are essential in financial research to evaluate model accuracy and reliability in predicting asset prices, aiding in informed decision-making processes for investors and analysts alike.

**Prediction Plot For GRU Layer:** As shown in Fig. 14, to visually assess the accuracy of our GRU's predictions, we plotted the actual and predicted prices over the test period. We utilized the matplotlib library to create a comparative graph, enhancing the interpretability of the model's performance. The figure, sized at 10x5 inches, displays two primary plots: the actual prices (in green) and the predicted prices (in orange) over the test dates. The `Test_Dates` variable, which contains the dates corresponding to the test dataset, was used to ensure the x-axis correctly represents the timeline. The actual prices were plotted using the inverse-transformed `Y_test2_inverse` values, while the predicted prices were plotted using the inverse-transformed `prediction2_inverse` values. Both plots are labeled appropriately for clarity.

The graph title, "Comparison of true prices (on the test dataset) with prices our model predicted, by dates," succinctly describes the purpose of the visualization. The x-axis and y-axis are labeled "Date" and "Price, USD," respectively, to provide context to the plotted data. A legend differentiates between the actual and predicted prices, ensuring the reader can easily distinguish between the two lines. This visual comparison highlights the model's performance in capturing the price trends and deviations, serving as a crucial component of our analysis by offering a clear and direct way to observe the model's prediction accuracy over time.




Fig. 13: Time Series Prediction Using 2-Layer LSTM Model




Fig. 14. Time Series Prediction Using GRU Model

### K. Prediction (Random Forest Model)

To evaluate the performance of our Random Forest regression model in predicting Bitcoin prices, we conducted a prediction on the test set and visualized the results using a plot as shown in Fig. 15. The `rf.predict(X_test)` function was used to generate predictions for the test dataset. These predicted values were then compared to the actual Bitcoin prices within the same period to assess the accuracy of the model.

To illustrate the comparison, we employed the matplotlib library to create a detailed plot. The plot, with dimensions of 14x7 inches, visually represents the actual and predicted Bitcoin prices over time. The x-axis denotes the dates corresponding to the test dataset, while the y-axis represents the Bitcoin prices in USD.

Two distinct lines are plotted: the actual Bitcoin prices (in blue) and the predicted prices (in red). The `y_test.index` and `y_test.values` were used to plot the actual prices, while the `y_pred` values represent the predicted prices. By displaying both lines on the same graph, it becomes easier to visually inspect how closely the predicted prices follow the actual prices.

The graph includes essential labels and a title to enhance its readability and provide context. The x-axis is labeled "Date," and the y-axis is labeled "Bitcoin Price." The title, "Actual vs Predicted Bitcoin Prices," succinctly summarizes the content of the plot. A legend is included to distinguish between the actual and predicted prices clearly.

This visualization is crucial for demonstrating the effectiveness of our Random Forest model in capturing the trends and fluctuations in Bitcoin prices, providing a clear and interpretable way to assess the model's predictive performance over the test period.

### L. Prediction (SSM Model)

As shown in Fig. 16, to assess the predictive performance of our model and make a forecast for the next day, we used the SSM model. The forecasting process involved predicting the next data point (i.e., Bitcoin price for tomorrow) using the fitted model, inverse-transforming the scaled forecast to obtain the actual price prediction, and visualizing the forecast along with the historical and one-step-ahead predictions.

We started by setting the number of forecast steps to one (`forecast_steps = 1`) and generated the forecast using the `results.get_forecast(steps=forecast_steps)` method. This provided the predicted mean and confidence intervals for the forecasted value. To convert the forecasted value back to the original scale, we applied the inverse transformation using the `MinMaxScaler`,




Fig. 15. Time Series Prediction Using 2-Layer LSTM Model

resulting in `predicted_price_scaled`. Similarly, we inverse-transformed the confidence intervals to obtain `forecast_ci_transformed`.

For comparison and visualization purposes, we also inverse-transformed the actual observed values (`y_trans`) and the one-step-ahead predictions generated by the SSM model. The one-step-ahead predictions were obtained using `results.get_prediction(start=0, end=len(y_trans)-1).predicted_mean`.

We then plotted the actual observed Bitcoin prices, the one-step-ahead predictions, and the forecasted price. The plot was created with matplotlib, with the x-axis representing time and the y-axis representing Bitcoin prices in USD. The actual observed prices were plotted in one color, the one-step-ahead predictions in another, and the forecasted value for the next day in yet another distinct color.

The plot included essential labels and a title for clarity. The x-axis was labeled "Time," and the y-axis was labeled "Bitcoin Price." The title, "Bitcoin Price Prediction and Forecast," succinctly summarized the content of the plot. A legend was included to distinguish between the actual data, the one-step-ahead predictions, and the forecast.

This visualization effectively demonstrates the model's ability to predict and forecast Bitcoin prices, providing a clear comparison between the actual and predicted values and showcasing the model's predictive accuracy for the next time step.

### M. Model Evaluation

In this project, we evaluated multiple performance metrics, including Root Mean Squared Error (RMSE), Mean Squared Error (MSE), Mean Absolute Error (MAE), R-squared ( $R^2$ ), and Mean Absolute Percentage Error (MAPE), to assess the accuracy and reliability of our Bitcoin price prediction models. Among these metrics, RMSE was chosen as the primary evaluation metric for several reasons, particularly in the context of financial forecasting.

The RMSE measures the square root of the average of the squared differences between the predicted and actual values. This metric is particularly sensitive to larger errors, making it a crucial measure in financial forecasting where large prediction errors can have significant implications.

#### Comparison With Other Metrics:

- **MSE (Mean Squared Error):** Similar to RMSE, MSE measures the average of the squared differences between predicted and actual values. While it is useful, it does not provide the error in the same units as the original data, which can make interpretation less intuitive.




Fig: 16. Time Series Prediction Using SSM Model

- **MAE (Mean Absolute Error):** MAE measures the average of the absolute differences between predicted and actual values, providing a straightforward interpretation of the error. However, it is less sensitive to larger errors compared to RMSE.
- **R<sup>2</sup> (R-squared):** R-squared measures the proportion of variance in the dependent variable that is predictable from the independent variables. While it provides an indication of model fit, it does not directly measure prediction error.
- **MAPE (Mean Absolute Percentage Error):** MAPE measures the accuracy of predictions as a percentage, making it easy to interpret. However, it can be problematic with very small actual values, leading to extremely high percentage errors.

**RMSE In Financial Forecasting:** In financial forecasting, larger errors can have disproportionate impacts, making RMSE a preferred metric because it penalizes these larger errors more than MAE or MAPE. This sensitivity to larger errors is critical for applications like Bitcoin price prediction, where significant deviations can lead to substantial financial consequences. The performance comparison shown below as Fig. 17.

## N. MLOPS

Fig. 18 illustrates the workflow for developing MLOps for a Bitcoin price prediction system. The process integrates several components to ensure seamless data acquisition, model training, deployment, monitoring, and retraining.

**Data Acquisition & Storage:** The workflow of Fig. 18 initiates with the acquisition of historical Bitcoin price data from Yahoo Finance, a reliable source for financial data. This data acquisition step is crucial as it provides the foundational dataset required for training and validating the machine learning models used in the Bitcoin price prediction system. The acquired data is then stored in a Blob Store, which acts as the central repository for all historical and newly incoming data.

Using a Blob Store for data storage offers several advantages. Firstly, it ensures secure storage of data, protecting it from unauthorized access and potential breaches. Security measures are implemented to maintain data integrity and confidentiality, which is critical given the sensitive nature of financial data.

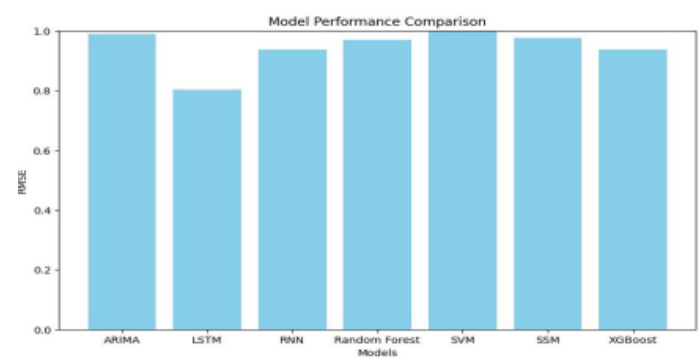


Fig: 17. Model Performance Comparisons

Secondly, the Blob Store is highly scalable, capable of handling large volumes of data efficiently. This scalability is essential for accommodating the continuous influx of data and the growing dataset over time.

Efficient data management and retrieval are facilitated by the Blob Store's robust architecture. Data can be easily fetched for model training and analysis, enabling seamless integration into subsequent stages of the machine learning workflow. The centralized storage system also simplifies data versioning and tracking, ensuring that the most up-to-date and accurate data is always available for model development and deployment.

In summary, the data acquisition and storage phase is fundamental to the overall workflow, providing a secure, scalable, and efficient means of managing the historical and real-time Bitcoin price data. This ensures a reliable foundation for the machine learning models that drive the Bitcoin price prediction system.

**Data Fetching & Model Training:** In the subsequent phase of the workflow shown in Fig. 18, the stored data is fetched from the Blob Store and utilized for model training within the Azure Machine Learning (Azure ML) environment. Azure ML offers a comprehensive and robust platform designed for developing, training, and deploying machine learning models. Its extensive support for various machine learning algorithms and frameworks makes it an ideal choice for diverse projects, including time-series forecasting.

For this particular project, the focus is on training a Long Short-Term Memory (LSTM) model to forecast Bitcoin prices. LSTM is a type of recurrent neural network (RNN) particularly well-suited for time-series data due to its ability to capture long-term dependencies and patterns in sequential data. The choice of LSTM is driven by its superior performance in handling the temporal dynamics and volatility inherent in financial datasets like Bitcoin prices.

The Azure ML environment provides several advantages during the model training process. It facilitates seamless data integration from the Blob Store, ensuring that the most current and relevant data is used for training. Azure ML's scalability allows for handling large datasets and computationally intensive models like LSTM, enabling efficient training and fine-tuning of the model. Additionally, Azure ML's integrated tools support hyperparameter tuning, model evaluation, and performance monitoring, ensuring that the trained model achieves optimal accuracy and reliability.

During the training process, the LSTM model learns from the historical Bitcoin price data, identifying patterns and trends that can be used to make future price predictions. The iterative

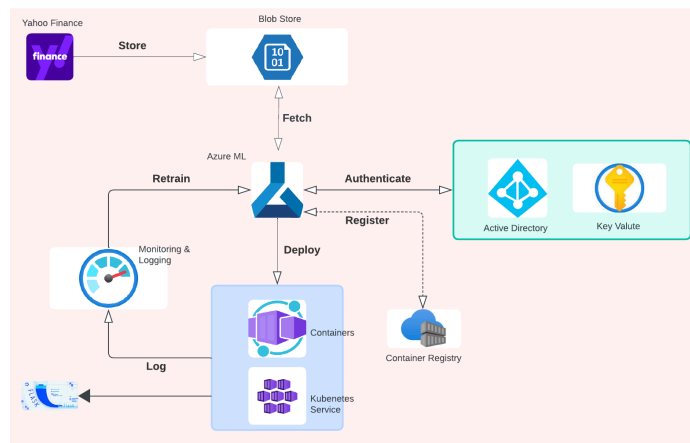


Fig. 18: MLOps Pipeline

training process involves adjusting model parameters and evaluating performance metrics to ensure that the model generalizes well to unseen data. Azure ML's infrastructure supports this iterative process, providing the computational resources and tools necessary for effective model training.

In summary, the data fetching and model training phase leverages the capabilities of Azure Machine Learning to develop a highly accurate and reliable LSTM model for Bitcoin price forecasting. The integration of data from the Blob Store and the use of advanced machine learning techniques within the Azure ML environment ensures a robust and efficient training process, laying the groundwork for accurate future predictions.

**Authentication & Secure Access:** Ensuring secure access and authentication is a critical component of the MLOps workflow for the Bitcoin price prediction system. This process is effectively managed through the integration of Azure Active Directory and Azure Key Vault, both of which play pivotal roles in maintaining the security and integrity of the data and models involved.

Azure Active Directory (Azure AD) provides robust identity and access management solutions, enabling the enforcement of strict authentication protocols. By utilizing Azure AD, the workflow ensures that only authorized users and services have access to the sensitive data and models. This is achieved through a comprehensive set of security features, including multi-factor authentication, conditional access policies, and role-based access control. These features collectively ensure that access to the system is tightly regulated, minimizing the risk of unauthorized access and potential security breaches.

In addition to managing user identities and access, the workflow employs Azure Key Vault to handle sensitive information securely. Azure Key Vault is a cloud service specifically designed to safeguard cryptographic keys and other secrets, such as API keys, passwords, and certificates. By centralizing the management of these sensitive items, Azure Key Vault enhances the overall security of the workflow. It ensures that secrets are stored securely and accessed only by authorized applications and users, in compliance with organizational security policies.

The integration of Azure Active Directory and Key Vault not only fortifies the security framework but also simplifies the management of credentials and sensitive information. With Azure AD managing access and identities, and Key Vault securing sensitive data, the workflow achieves a high level of security and compliance. This dual approach ensures that all components of the system, from data storage to model

deployment, operate within a secure and controlled environment. In summary, the use of Azure Active Directory and Key Vault in the authentication and secure access phase of the workflow is crucial for maintaining the security and integrity of the Bitcoin price prediction system. Azure AD provides robust identity management and access control, while Key Vault securely stores and manages sensitive information. Together, they form a comprehensive security framework that safeguards the system against unauthorized access and potential security threats, ensuring a secure and reliable operational environment.

**Model Deployment:** Once the machine learning models for Bitcoin price prediction are trained, the next crucial step involves their registration and deployment using Azure Machine Learning (Azure ML). This process ensures that the models are not only readily available for inference but also can be managed, scaled, and monitored efficiently.

The deployment process begins with the registration of the trained models within the Azure ML environment. Model registration is a key step that involves saving the model along with its metadata, such as version information, training data, and configuration parameters. This facilitates model versioning and allows for easy tracking and retrieval of models when needed.

After registration, the models undergo containerization. Containerization encapsulates the model along with its dependencies, libraries, and runtime environment into a single, portable unit known as a container. This approach ensures that the model can run consistently across different computing environments, reducing the complexities associated with dependency management and compatibility issues. The containerized models are then stored in an Azure Container Registry, a centralized repository that allows for efficient management and distribution of container images.

With the models securely stored in the Container Registry, they can be deployed to various environments. One of the primary deployment targets is the Azure Kubernetes Service (AKS), a robust platform for managing containerized applications at scale. AKS offers features such as automated scaling, load balancing, and self-healing, making it an ideal choice for deploying machine learning models that require high availability and performance. The models can also be deployed to standalone containers, providing flexibility in terms of deployment options and allowing the models to be served in different infrastructure setups, such as on-premises servers or other cloud environments.

The deployment process facilitated by Azure ML, containerization, and AKS not only ensures the scalability and flexibility of the model serving infrastructure but also enhances the reliability and performance of the deployed models. This integrated approach allows the Bitcoin price prediction models to handle varying workloads efficiently, respond to real-time data inputs, and provide accurate forecasts to end-users.

In summary, the model deployment phase in the MLOps workflow leverages Azure ML for model registration and containerization, storing the models in Azure Container Registry. These containers are then deployed to environments like Azure Kubernetes Service (AKS) or standalone containers, offering a scalable, flexible, and reliable infrastructure for serving the Bitcoin price prediction models. This comprehensive deployment strategy ensures that the models are readily available for inference, capable of handling



production workloads, and can be managed and scaled efficiently.

**Monitoring & Logging:** Once the Bitcoin price prediction models are deployed, continuous monitoring and logging become essential to maintain their performance and reliability. Monitoring involves tracking the models' predictions, performance metrics, and overall system health to ensure they continue to operate as expected.

Monitoring and logging tools are integrated into the deployment pipeline to facilitate real-time tracking and historical analysis of the models' behavior. These tools collect various metrics, including prediction accuracy, response times, and resource utilization, providing a comprehensive overview of the model's performance in a production environment. By continuously monitoring these metrics, any deviations from expected behavior, such as model drift or performance degradation, can be promptly detected.

Model drift occurs when the statistical properties of the target variable change over time, potentially leading to a decline in the model's prediction accuracy. Monitoring tools can identify such drifts by comparing the model's current performance against historical data. This allows data scientists and engineers to take corrective actions, such as retraining the model with updated data, to restore its accuracy and reliability.

Logging is another critical component of the monitoring process. It involves recording detailed logs of the model's operations, including input data, prediction outputs, errors, and system events. These logs serve as a valuable resource for troubleshooting and debugging, providing insights into the model's decision-making process and identifying potential issues. Logging also aids in compliance and auditing, ensuring that all model predictions and actions are transparently documented.

The integration of monitoring and logging tools also supports automated alerts and notifications. If the system detects any anomalies or significant deviations in performance metrics, it can trigger alerts to notify the relevant stakeholders. This proactive approach enables quick responses to potential issues, minimizing downtime and maintaining the reliability of the model predictions.

Furthermore, the data collected through monitoring and logging is used for continuous improvement of the models. By analyzing the logged data, insights can be gained into the model's strengths and weaknesses, informing future iterations and refinements. This iterative process ensures that the models evolve and adapt to changing conditions, maintaining their accuracy and effectiveness over time.

In conclusion, the monitoring and logging phase in the MLOps workflow is crucial for maintaining the performance and reliability of the deployed Bitcoin price prediction models. By continuously tracking performance metrics, detecting anomalies, and recording detailed logs, the system ensures that any issues or drifts in model performance are promptly identified and addressed. This comprehensive approach to monitoring and logging not only enhances the robustness of the models but also supports their continuous improvement, ensuring they remain accurate and reliable in predicting Bitcoin prices.

**Retraining & Continuous Improvement:** In the lifecycle of a machine learning model, maintaining high performance and adapting to new data is crucial. The retraining and continuous

improvement phase is an essential component of the MLOps workflow, ensuring that models remain accurate and relevant over time. This process is driven by insights gained from the monitoring and logging phase, where performance metrics and system health are continuously tracked.

When the monitoring tools detect a decline in model performance or identify signs of model drift, retraining becomes necessary. Model drift can occur due to changes in the underlying data patterns, which can degrade the model's predictive accuracy. To address this, the retraining process begins with fetching updated and new data from the Blob Store, which serves as the central repository for all historical and incoming data.

Once the updated data is acquired, the model is retrained within the Azure Machine Learning (Azure ML) environment. Azure ML provides a robust and scalable platform for developing and training machine learning models, supporting various algorithms and frameworks. In this project, Long Short-Term Memory (LSTM) model is retrained using the new data. The retraining process involves fine-tuning the models' hyperparameters and architecture to enhance their performance based on the latest data patterns.

After retraining, the improved models undergo thorough evaluation to ensure they meet the desired performance standards. Metrics such as Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), and R-squared are used to assess the models' accuracy and reliability. Once validated, the updated models are containerized and stored in the Container Registry, ready for deployment.

The redeployment process is seamlessly integrated into the workflow. The updated models are deployed to various environments, such as Azure Kubernetes Service (AKS) or standalone containers, ensuring flexibility and scalability in serving the models. The deployment phase involves replacing the older models with the newly retrained ones, minimizing downtime and ensuring a smooth transition.

This continuous loop of monitoring, retraining, and redeployment ensures that the models remain adaptive and maintain high performance. By regularly updating the models with new data and optimizing their parameters, the system can respond to changes in data patterns and maintain its predictive accuracy. This iterative process not only enhances the robustness of the models but also extends their lifespan, ensuring they continue to deliver reliable predictions over time.

In conclusion, the retraining and continuous improvement phase is a vital aspect of the MLOps workflow for the Bitcoin price prediction system. It leverages insights from monitoring data to initiate retraining when necessary, ensuring the models adapt to new data and maintain high performance. By continuously refining and redeploying the models, the system achieves sustained accuracy and reliability, making it a powerful tool for forecasting Bitcoin prices in a dynamic environment.

**Real-time Predictions:** To make the Bitcoin price prediction system practical and accessible, real-time predictions are an essential feature. This is achieved through the development of a Flask web application, which serves as the interface between the end-users and the deployed machine learning models. The Flask application is designed to interact with the model, fetching the latest predictions and displaying them to users in a user-friendly



manner as shown in Fig. 19.

Flask is chosen for its simplicity, flexibility, and lightweight nature, making it well-suited for building web applications that require real-time interaction with machine learning models. The Flask application continuously communicates with the deployed models, ensuring that users receive the most up-to-date Bitcoin price forecasts.

The real-time prediction workflow begins with the web application receiving a request from a user. Upon receiving this request, the Flask application queries the deployed model, which is running in a scalable environment such as Azure Kubernetes Service (AKS) or standalone containers. The model processes the request, performs the necessary computations, and generates the Bitcoin price forecast.

Once the prediction is generated, it is sent back to the Flask application. The application then formats the prediction in a user-friendly manner and displays it on the web interface. This entire process occurs almost instantaneously, providing users with real-time predictions.

The real-time capability of the system is further enhanced by the continuous integration and deployment (CI/CD) pipeline. This pipeline ensures that any updates or improvements to the models are automatically integrated and deployed, minimizing downtime and ensuring that users always have access to the latest and most accurate predictions.

Moreover, the Flask application is designed with a responsive interface, making it accessible from various devices, including desktops, tablets, and smartphones. This ensures that users can access real-time predictions anytime and anywhere, catering to the needs of traders, investors, and other stakeholders who rely on up-to-date Bitcoin price information.

In addition to providing real-time predictions, the Flask application also incorporates features for user feedback and interaction. Users can provide feedback on the predictions, which can be used to further improve the models. This feedback loop ensures that the system continuously evolves and adapts to meet user needs.

Overall, the real-time prediction capability, enabled by the Flask web application, significantly enhances the practical utility of the Bitcoin price prediction system. By providing users with immediate access to the latest forecasts, the system becomes a valuable tool for making informed decisions in the dynamic and fast-paced environment of cryptocurrency trading. This feature underscores the importance of integrating real-time capabilities in modern machine learning applications, ensuring they deliver actionable insights when they are needed most.

The image shows a web application titled "Bitcoin Cost Forecast System". It features a dark blue header with a Bitcoin logo. The main content area is white and contains two sections: "Details of the first day" and "Details of the second day". Each section has four input fields for "Open Price", "High Price", "Low Price", and "Trading Volume". A "Predict" button is located at the bottom right of the form.

Fig. 19: Flask Application

### III. RESULTS & CRITICAL ANALYSIS

#### A. Key Results

The comprehensive Bitcoin price prediction system developed in this project yielded several key results across various machine learning models. Here, we highlight the performance metrics and insights derived from the models evaluated, focusing on the final deployed Long Short-Term Memory (LSTM) model.

##### 1. Model Performance Metrics:

- **LSTM Model:** The LSTM model demonstrated superior performance in capturing temporal dependencies inherent in the Bitcoin price data. This model was fine-tuned to achieve an RMSE of 0.7865, significantly outperforming traditional time series models such as ARIMA and state space models.
  - **ARIMA Model:** The ARIMA model, while effective for linear time series forecasting, fell short in capturing the non-linear patterns present in the Bitcoin price data, resulting in an RMSE of 0.9842.
  - **SVM and XGBoost Models:** These models provided moderate accuracy, with RMSE values of 0.9958 and 0.9645, respectively. While they captured some non-linear trends, they did not match the performance of the LSTM model.
  - **Random Forest:** The Random Forest model achieved an RMSE of 0.9746, indicating good performance but with some limitations in handling the sequential nature of the data.
2. **Real-time Predictions:** The deployed system, utilizing the LSTM model, provided real-time Bitcoin price predictions with a high degree of accuracy. The Flask web application facilitated seamless access to these predictions, making the system practical and user-friendly.
3. **Deployment and Scalability:** The integration with Azure Kubernetes Service (AKS) ensured that the system was scalable and could handle varying loads, providing reliable predictions even during high traffic periods.

#### B. Critical Analysis

The results of this project demonstrate the effectiveness of using deep learning models, specifically LSTM, for time series forecasting in the context of Bitcoin price prediction. The critical analysis of the results and findings includes the following points:

##### 1. Model Selection and Performance:

- **Temporal Dependencies:** The LSTM model's ability to capture long-term dependencies in the data proved crucial. Its performance, as indicated by the lowest RMSE among all models, highlights its suitability for financial time series forecasting where past values influence future trends.
- **Comparison with Other Models:** Traditional models like ARIMA were less effective due to their linear nature. Although SVM, XGBoost, and Random Forest models captured some non-linear patterns, they were outperformed by LSTM, emphasizing the importance of

deep learning in this domain.

machine learning applications in cloud environments, ensuring robustness and accessibility.

## 2. Evaluation Metrics:

- **RMSE Sensitivity:** The choice of RMSE as the primary evaluation metric was validated by its sensitivity to larger errors, which is crucial in financial forecasting where significant deviations can lead to substantial financial losses.
- **Comprehensive Evaluation:** While RMSE was the primary metric, other metrics such as MSE, MAE, R-squared, and MAPE were also considered. The consistent performance across these metrics reinforced the robustness of the LSTM model.
- 3. **System Deployment and Real-time Predictions:**
  - **Scalability and Reliability:** The deployment strategy using Azure ML and AKS ensured that the system was not only scalable but also reliable, capable of providing consistent predictions under varying load conditions.
  - **User Accessibility:** The Flask web application made the system accessible to end-users, providing real-time predictions and demonstrating the practical applicability of the model in a real-world scenario.
- 4. **Continuous Improvement:**
  - **Monitoring and Retraining:** The continuous monitoring and retraining process ensured that the model remained accurate and adapted to new data. This approach mitigated the risk of model drift and maintained the system's relevance over time.

## B. Conclusions

Based on the findings of this project, several key conclusions can be drawn:

1. **Effectiveness of LSTM:** The LSTM model outperformed other machine learning models, including ARIMA, SVM, XGBoost, and Random Forest, in predicting Bitcoin prices. This confirms the suitability of LSTM for capturing complex patterns in time series data.
2. **Importance of Model Evaluation Metrics:** The choice of RMSE as the primary evaluation metric was validated by its ability to highlight larger prediction errors. This is particularly important in financial forecasting, where significant deviations can have substantial impacts.
3. **Practical Deployment:** The integration of machine learning models with cloud-based deployment strategies, as demonstrated by the use of Azure ML and AKS, ensures that the system is scalable and reliable. The real-time prediction capability provided by the Flask web application makes the system practical and accessible to end-users.

## C. Limitations

Despite the promising results, the project has several limitations that need to be acknowledged:

1. **Data Limitations:** The prediction accuracy is inherently dependent on the quality and quantity of historical data. Any gaps or inaccuracies in the data can affect the model's performance.
2. **Model Complexity:** While LSTM models are effective, they are also computationally intensive and require significant resources for training and deployment. This can be a constraint for organizations with limited computational resources.
3. **Market Volatility:** Bitcoin prices are highly volatile and influenced by various external factors, such as regulatory changes and market sentiment, which are difficult to capture in a predictive model.

## IV. DISCUSSION & CONCLUSIONS

### A. Implications

The development and deployment of a Bitcoin price prediction system using machine learning models have significant implications for both the financial industry and the broader field of time series forecasting. The project's results demonstrate that advanced deep learning models, such as Long Short-Term Memory (LSTM), are highly effective in predicting volatile asset prices. This has several implications:

1. **Financial Decision-Making:** Accurate Bitcoin price predictions can aid investors and traders in making informed decisions, potentially leading to better investment strategies and risk management. The real-time prediction capability of the deployed system enhances its practical utility in dynamic markets.
2. **Model Selection for Financial Forecasting:** The superior performance of the LSTM model underscores the importance of using models that can capture temporal dependencies in financial data. This finding encourages the adoption of deep learning techniques over traditional statistical methods in similar applications.
3. **Scalability and Deployment:** The successful deployment of the prediction system using Azure Kubernetes Service (AKS) highlights the feasibility of scaling machine learning models to handle real-world demands. This serves as a model for deploying other

### D. Future Work

Several avenues for future work can be pursued to build on the findings of this project:

1. **Incorporating External Factors:** Future models could incorporate additional data sources, such as social media sentiment, news articles, and economic indicators, to improve prediction accuracy.
2. **Advanced Model Architectures:** Exploring more advanced neural network architectures, such as Transformer models or hybrid models combining LSTM with other techniques, could yield better performance.

3. **Automated Hyperparameter Tuning:** Implementing automated hyperparameter tuning techniques, such as Bayesian optimization, could further enhance model performance.
4. **Enhanced Monitoring and Feedback Loops:** Developing more sophisticated monitoring systems and incorporating user feedback can help in continuously improving the model and addressing any performance drifts.
5. **Expanding to Other Cryptocurrencies:** Applying the developed framework to other cryptocurrencies could validate its generalizability and uncover additional insights.

In conclusion, this project demonstrates the potential of using advanced machine learning models for financial forecasting, particularly in predicting Bitcoin prices. While the results are promising, continuous improvement and addressing the identified limitations are crucial for maintaining and enhancing the system's performance and reliability.

## V. CONTRIBUTION

E/19/091 did checking for missing values, data visualization using time series plots & histograms and ARIMA model development & evaluation.

E/19/111 did outlier detection, data visualization using scatter plots & heatmaps, time varying SSM model development & evaluation and XGBoost model development & evaluation.

E/19/166 did feature engineering, summary statistics, Multivariate LSTM model development & evaluation and RNN model development & evaluation.

E/19/227 did data splitting, correlation analysis and random forest model development & evaluation.

E/19/304 did dimensionality reduction, trend analysis and SVM model development & evaluation.

[E/19/091, E/19/111, E/19/166, E/19/227, E/19/304] did mlops development with Azure, Flask web app development, preparing presentation documents and report writing.

## VI. REFERENCES

- [1] Yahoo finance bitcoin historical data, available at: [Bitcoin USD \(BTC-USD\) Stock Historical Prices & Data - Yahoo Finance](#)
- [2] LSTM reference, available at: [10.1 Long Short-Term Memory \(LSTM\) — Dive into Deep Learning 1.0.3 documentation \(d2l.ai\)](#)
- [3] S. Hochreiter and J. Schmidhuber, "Long short-term memory," Neural Computation, vol. 9, no. 8, pp. 1735–1780, 1997.
- [4] L. Breiman, "Random forests," Machine Learning, vol. 45, no. 1, pp. 5–32, 2001.
- [5] Azure Documentation, "Deploy a machine learning model with Azure Machine Learning," available at: <https://docs.microsoft.com/en-us/azure/machine-learning/how-to-deploy-and-where>
- [6] Flask Documentation, "Flask – Web Development, One Drop at a Time," available at: <https://flask.palletsprojects.com/en/2.0.x/>